

Process Synchronization

Dr. SUBHASIS DASH
SCHOLE OF COMPUTER ENGINEERING.
KIIT UNIVERSITY
BHUBANESWAR

Cooperating Processes

- ❑ **Independent** process cannot affect or be affected by the execution of another process
- ❑ **Cooperating** process can affect or be affected by the execution of another process
- ❑ Advantages of process cooperation
 - ❑ Information sharing
 - ❑ Computation speed-up
 - ❑ Modularity
 - ❑ Convenience

Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
 - *unbounded-buffer* places no practical limit on the size of the buffer
 - *bounded-buffer* assumes that there is a fixed buffer size

Bounded-Buffer – Shared-Memory Solution

□ Shared data

```
#define BUFFER_SIZE    n
typedef struct {
    . . .
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

□ Solution is correct, but can only use BUFFER_SIZE is [n-1] elements.

- `((in = (in + 1) % BUFFER_SIZE count) == out) :- buffer full`
- `(in == out) :- buffer empty`

Background

- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers. We can do so by having an integer **count** that keeps track of the number of full buffers. Initially, count is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.

Producer

```
while (true) {  
  
    /* produce an item and put in nextProduced */  
    while (count == BUFFER_SIZE)  
        ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

Consumer

```
while (true) {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
  
    /* consume the item in nextConsumed  
  
    */  
}
```

Producer & Consumer Code

```
while (true)
{
    /* produce an item and put in next
    Produced */
    while (count == BUFFER_SIZE);
    // do nothing
    buffer [in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
    count++;
}
```

```
while (true)
{
    while (count == 0);
    // do nothing
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    count--;
    /* consume the item in next
    Consumed */
}
```

put side by side

Race Condition

- `count++` could be implemented as

```
register1 = count  
register1 = register1 + 1  
count = register1
```

- `count--` could be implemented as

```
register2 = count  
register2 = register2 - 1  
count = register2
```

- Consider this execution interleaving with “count = 5” initially:

```
S0: producer execute register1 = count {register1 = 5}  
S1: producer execute register1 = register1 + 1 {register1 = 6}  
S2: consumer execute register2 = count {register2 = 5}  
S3: consumer execute register2 = register2 - 1 {register2 = 4}  
S4: producer execute count = register1 {count = 6}  
S5: consumer execute count = register2 {count = 4}
```

CRITICAL SECTION (CS) PROBLEM

- ❑ Critical is a segment of code in each process where process wants to change there common variable.
- ❑ But CS allows only one process to execute in its CS –: **MUTUALLY EXCLUSIVE.**
- ❑ There are 3 different sections:-
 - ❑ ENTRY SECTION
 - ❑ EXIT SECTION
 - ❑ REMINDER SECTION

Example:-

repeat

ENTRY

critical section

EXIT

reminder section

until false

Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then only those process that are not executing their reminder section can participate in the selection of the processes that which will enter the critical section next, and this selection can't be postponed indefinitely
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted .
 - Assume that each process executes at a nonzero speed
 - No assumption concerning relative speed of the N processes


1st S/W Solution for CS problem

□ P_i 's Algorithm

repeat

var turn = i

While (turn \neq " i ") do no-op;



Critical section

turn=" j ";


until false

□ P_j 's Algorithm

repeat

var turn = j

While (turn \neq " j ") do no-op;



Critical section

turn=" i ";

until false


2nd Solution

□ P_i 's Algorithm

Var flag : array [i...j] of boolean
repeat

flag [i] = true;

While (flag [j] == " true ") do no-op;



Critical section

Flag [i]=" false ";


until false

□ P_j 's Algorithm

Var flag : array [i...j] of boolean
repeat

flag [j] = true;

While (flag [i] == " true ") do no-op;



Critical section

Flag [j]=" false ";

until false

Peterson's Solution

- Two process solution
- Assume that the LOAD and STORE instructions are atomic; that is, cannot be interrupted.
- The two processes share two variables:
 - int **turn**;
 - Boolean **flag[2]**
- The variable **turn** indicates whose turn it is to enter the critical section.
- The **flag** array is used to indicate if a process is ready to enter the critical section. **flag[i]** = true implies that process P_i is ready!

Peterson's Solution

□ P_i 's Algorithm

repeat

S1: flag [i] = TRUE;

S2: turn = j;

S3: while (flag[j] == TRUE &&
turn != i) do no-op;

CRITICAL SECTION

flag [i] = FALSE;

until false

□ P_j 's Algorithm

repeat

L1: flag [j] = TRUE;

L2: turn = i;

L3: while (flag [i] == TRUE &&
turn != j) do no-op;

CRITICAL SECTION

flag [j] = FALSE;

until false

Sequence of execution:

S1; S2; L1;L2;

S1;L1;L2;S2;

L1;S1;L2;S2

Synchronization Hardware

- Many systems provide hardware support for critical section code
- Uniprocessors – could disable interrupts
 - Currently running code would execute without preemption
 - Generally too inefficient on multiprocessor systems
 - ▶ Operating systems using this not broadly scalable
- Modern machines provide special atomic hardware instructions
 - ▶ **Atomic = non-interruptable**
 - Either test memory word and set value
 - Or swap contents of two memory words

Solution to Critical-section Problem Using Locks

```
do {  
    acquire lock  
        critical section  
    release lock  
        remainder section  
} while (TRUE);
```

→ A LOCK is a object that provides the following 2 operations! -

① acquire () → wait to enter CS

② release () → Allow another to enter to CS.

acquire_lock ()

Critical Section

Release_lock ()

Remainder Section

Ex: int withdraw(account, amount)
{
 acquire(lock);
 balance = get_balance(account);
 balance = balance - amount;
 Put_balance(account, balance);
 release(lock);
}

Return balance;

Lock implementation

Struct lock

```
{  
    int held = 0;  
}
```

void acquire(lock)

```
{  
    while (lock->held);  
    lock->held = 1;  
}
```

void release(lock)

```
{  
    lock->held = 0;  
}
```

→ it inst.ⁿ preempted after while then mutual exclusion violated. // To prevent it disable interrupt.
→ Busy waits
→ context switch here.
→ Hence this sequence must be atomic. //

↳ one solution.

TestAndndSet Instruction

□ Definition:

```
boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

Solution using TestAndSet

- Shared boolean variable lock., initialized to false.
- Solution:

```
while (true) {  
    while ( TestAndSet (&lock ))  
        ; /* do nothing  
  
        // critical section  
  
    lock = FALSE;  
  
    // remainder section  
  
}
```

Swap Instruction

□ Definition:

```
void Swap (boolean *a, boolean *b)
{
    boolean temp = *a;
    *a = *b;
    *b = temp;
}
```

Solution using Swap

- ❑ Shared Boolean variable lock initialized to FALSE; Each process has a local Boolean variable key.

- ❑ Solution:

```
while (true) {  
    key = TRUE;  
    while ( key == TRUE)  
        Swap (&lock, &key );  
  
        //  critical section  
  
    lock = FALSE;  
  
        //  remainder section  
  
}
```

Semaphore

- Synchronization tool that does not require busy waiting
- Semaphore (S) is an integer variable in which two atomic operations P & V is defined to provide synchronization & mutual exclusion in concurrent system.
- Two standard operations modify S: `wait()` and `signal()`
 - Originally called `P(S)` “TO TEST” and `V(S)` “TO INCRIMENT”
- Less complicated
- Can only be accessed via two indivisible (atomic) operations
 - `P(S) : wait (S)`

```
{  
    while S <= 0; // no-op  
    S--;  
}
```
 - `V(S) : Signal (S)`

```
{  
    S++;  
}
```


Semaphore as General Synchronization Tool

- **Counting** semaphore – integer value can range over an unrestricted domain
- **Binary** semaphore – integer value can range only between 0 and 1; can be simpler to implement
 - Also known as **mutexlocks**
- Can implement a counting semaphore **S** as a binary semaphore
- Provides mutual exclusion
 - Semaphore **S**; // initialized to 1
 - wait (**S**);
Critical Section
signal (**S**);

Semaphore Implementation

- ❑ Must guarantee that no two processes can execute `wait ()` and `signal ()` on the same semaphore at the same time
- ❑ Thus, implementation becomes critical due to critical section problem where the wait and signal code both are placed in the critical section.
 - ❑ Could now have **busy waiting** in critical section implementation *[which demands mutual exclusion]*
 - ▶ ***Spin Lock***
 - ▶ But implementation code is short
 - ▶ Little busy waiting if critical section rarely occupied
- ❑ Note that applications may spend lots of time in critical sections and therefore this is not a good solution.

Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue. Each entry in a waiting queue has two data items:
 - value (of type integer)
 - pointer to next record in the list

- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue.
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue.

Semaphore Implementation with no Busy waiting (Cont.)

- Implementation of wait:

```
wait (S){  
    value--;  
    if (value < 0) {  
        add this process to waiting queue  
        block(); }  
}
```

- Implementation of signal:

```
Signal (S){  
    value++;  
    if (value <= 0) {  
        remove process P from waiting queue  
        wakeup(P); }  
}
```

Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let **S** and **Q** be two semaphores initialized to 1

P_0	P_1
wait (S);	wait (Q);
wait (Q);	wait (S);
.	.
.	.
.	.
signal (S);	signal (Q);
signal (Q);	signal (S);

- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.

Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem

Bounded-Buffer Problem

- N buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
- Semaphore **full** initialized to the value 0
- Semaphore **empty** initialized to the value N .

Bounded Buffer Problem (Cont.)

- N buffers, each can hold one item
- Semaphore **mutex** = 1
- Semaphore **full** = 0
- Semaphore **empty** = N .

Structure of producer process

```
while (true) {  
    // produce an item  
    wait (empty);  
    wait (mutex);  
    // add the item to  
    the buffer  
    signal (mutex);  
    signal (full);  
}
```

Structure of consumer process

```
while (true) {  
    wait (full);  
    wait (mutex);  
    // remove an item  
    from buffer  
    signal (mutex);  
    signal (empty);  
    // consume the  
    removed item  
}
```


Readers-Writers Problem

- A data set is shared among a number of concurrent processes
 - Readers – only read the data set; they do **not** perform any updates
 - Writers – can both read and write.
- Problem – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time.
- Shared Data
 - Data set
 - Semaphore **mutex** initialized to 1.
 - Semaphore **wrt** initialized to 1.
 - Integer **readcount** initialized to 0.

Readers-Writers Problem (Cont.)

- Structure of writer process

```
while (true)
{
    wait (wrt) ;

    //writing is
    performed


    signal (wrt) ;
}
```

- Structure of reader process

```
while (true)
{
    wait (mutex) ;
    readcount ++ ;
    if (readcount == 1) wait (wrt) ;
    signal (mutex)

    // reading is performed
    wait (mutex) ;
    readcount - - ;
    if (readcount == 0) signal (wrt) ;
    signal (mutex)
}
```

Dining-Philosophers Problem



□ Shared data

- Bowl of rice (data set)
- Semaphore **chopstick** [5] initialized to 1

Dining-Philosophers Problem (Cont.)

- The structure of Philosopher *i*:

```
While (true)
```

```
{
```

```
    wait ( chopstick[i] );
```

```
    wait ( chopStick[ (i + 1) % 5] );
```

```
        // eat
```

```
    signal ( chopstick[i] );
```

```
    signal ( chopstick[ (i + 1) % 5] );
```

```
        // think
```

```
}
```

Problems with Semaphores

- ❑ Incorrect use of semaphore operations:
 - ❑ `signal (mutex) wait (mutex)`
 - ❑ `wait (mutex) ... wait (mutex)`
 - ❑ Omitting of `wait (mutex)` or `signal (mutex)` (or both)

Monitors


- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- Only one process may be active within the monitor at a time

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }
    ...

    procedure Pn (...) {.....}

    Initialization code ( ....) { ... }
    ...
}
}
```

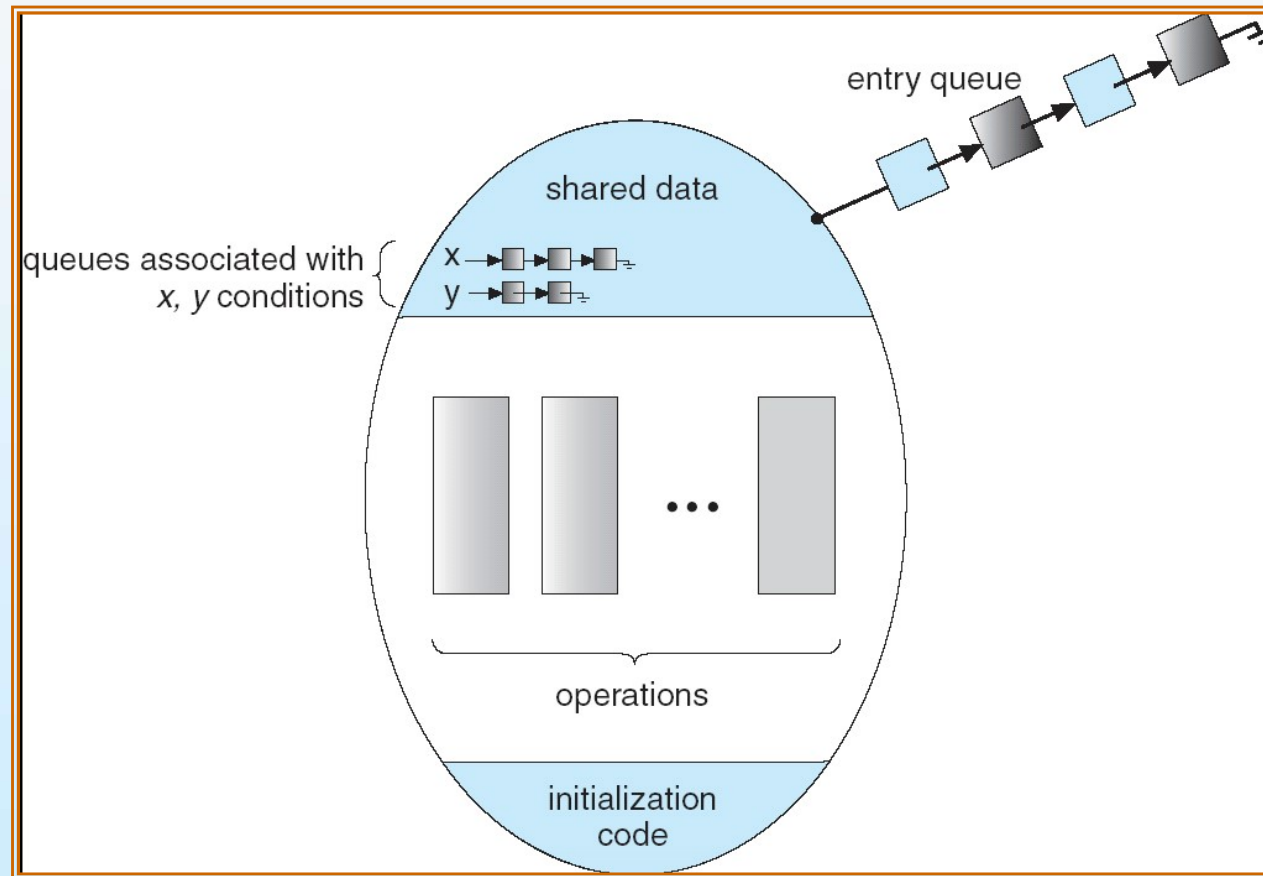
Schematic view of a Monitor



Condition Variables

- `condition x, y;`
- Two operations on a condition variable:
 - `x.wait ()` – a process that invokes the operation is suspended.
 - `x.signal ()` – resumes one of processes (if any) that invoked `x.wait ()`

Monitor with Condition Variables



- If process P invokes `x.signal()`, and process Q is suspended in `x.wait()`, what should happen next?
 - Both Q and P cannot execute in parallel. If Q is resumed, then P must wait
- Options include
 - **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition
 - **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition
 - Both have pros and cons – language implementer can decide
 - Monitors implemented in Concurrent Pascal compromise
 - ▶ P executing signal immediately leaves the monitor, Q is resumed
 - Implemented in other languages including Mesa, C#, Java

Monitor Solution to Dining Philosophers

```
monitor DiningPhilosophers
{
    enum { THINKING; HUNGRY, EATING) state [5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING) self[i].wait;
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```

Monitor Solution to Dining Philosophers

```
void test (int i) {  
    if ((state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) ) {  
        state[i] = EATING ;  
        self[i].signal () ;  
    }  
}
```

```
    initialization_code() {  
        for (int i = 0; i < 5; i++)  
            state[i] = THINKING;  
    }  
}
```

Monitor Solution to Dining Philosophers

- Each philosopher i invokes the operations `pickup()` and `putdown()` in the following sequence:

`DiningPhilosophers.pickup(i);`

EAT

`DiningPhilosophers.putdown(i);`

- No deadlock, but starvation is possible

Bounded buffer Solution using monitor

– Bounded buffer using monitors and signals

- * Shared State data[10] - a buffer holding produced data.
num - tells how many produced data items there are in the buffer.
- * Atomic Operations Produce(v) called when producer produces data item v.
Consume(v) called when consumer is ready to consume a data item. Consumed item put into v.
- * Condition Variables bufferAvail - signalled when a buffer becomes available.
dataAvail - signalled when data becomes available.

```
monitor PC {  
    Condition *bufferAvail, *dataAvail;  
    int num = 0;  
    int data[10];  
  
    Produce(v) {  
        while (num == 10) {  
            bufferAvail→Wait();  
        }  
        put v into data array  
        num++;  
        dataAvail→Signal();  
    }  
    Consume(v) {  
        while (num == 0) {  
            dataAvail→Wait();  
        }  
        put next data array value into v  
        num-;  
        bufferAvail→Signal();  
    }  
}
```

End of Chapter 4